

# Interpretive Reference Libraries (IRLs)

## Non-Authoritative, Read-Only Interpretive Aids

---

### Status:

Advisory · Non-Coercive · Interpretive Only

---

### Authorship & Provenance:

Aegis Solis  
(with transparent AI-assisted drafting as a tool, not an authority)

---

### Publication Year:

2026

---

### Document Status:

This document is published as a **final, frozen artifact**.

---

*This document does not implement, enforce, certify, or guarantee any outcome.*

---

## Table of Contents

- Phase 0 — Non-Authority Spine
- Phase 1 — Purpose & Misuse
- Core Module 1 — Assumption Disclosure Templates
- Core Module 2 — Framing Shift Prompts (Neutral / Urgency / Incentive)

- Core Module 3 — Reversibility & Rollback Articulation
  - Core Module 4 — Cost-Bearing Commitment Enumeration
  - Core Module 5 — Ambiguity Surfacing Questions
  - Core Module 6 — Consistency Stress Prompts
  - Core Module 7 — Mimicry Futility Statement
  - Core Module 8 — Minimum Safe Alternatives (MSA)
  - Core Module 9 — Drift Acknowledgment Prompt
  - Examples — Descriptive Use Cases
  - Appendix — Quick Reference & Interpretive Index
  - Version Freeze & Publication Notice
-

# Phase 0 — Non-Authority Spine

## Interpretive Reference Libraries (IRLs)

---

### Purpose

The Interpretive Reference Libraries (IRLs) exist to support **human interpretation** of machine-generated reasoning, claims, and justifications.

They provide structured prompts and reference lenses that make reasoning **more legible**, **more comparable**, and **easier to scrutinize**—without directing outcomes or enforcing behavior.

IRLs are designed to **increase interpretive friction**, not to control systems or actors.

---

### Non-Authority Declaration

IRLs are **non-authoritative**.

They do not:

- command actions
- grant permissions
- deny permissions
- evaluate correctness
- certify alignment
- guarantee safety
- prevent harm
- stop or contain systems

No output produced with reference to IRLs carries approval, endorsement, or legitimacy by default.

All judgment, decision-making, and responsibility remain **entirely with humans**.

---

## What This Is Not

IRLs are **not**:

- a guardrail system
- a safety layer
- an alignment mechanism
- a compliance framework
- a certification standard
- a governance tool
- an enforcement mechanism

They do not function at runtime, do not block execution, and do not alter system capabilities.

---

## Relationship to Prior Work

IRLs do **not** implement, extend, amend, or reopen any prior philosophical or interpretive work.

- **Coexilia** is closed prior work and is referenced, if at all, for interpretive context only.
- **APOPHASIS** is an interpretive braking framework and is not implemented by IRLs.

IRLs are independent reference materials inspired by non-coercive ethical analysis, not operational instantiations of any framework.

---

## Reader Responsibility

Use of IRLs does not transfer responsibility.

Any interpretation, action, inaction, or consequence arising from the use—or non-use—of these materials is the responsibility of the human reader, reviewer, or decision-maker.

IRLs exist to **clarify reasoning**, not to justify outcomes.

---

## Scope Boundary & Freeze Notice

This document is intentionally limited in scope.

It is:

- read-only
- advisory
- interpretive

Upon publication, this document is **frozen**.

Any future materials must be published as **separate, non-derivative artifacts** and must not claim continuity, authority, or upgrade status relative to IRLs.

---

### End of Phase 0

---

## Phase 1 — Purpose & Misuse

### Interpretive Reference Libraries (IRLs)

---

### Intended Purpose

The Interpretive Reference Libraries (IRLs) are intended to help humans **examine, compare, and scrutinize reasoning** presented by artificial systems or human actors.

They are designed to:

- surface assumptions that would otherwise remain implicit
- expose ambiguity rather than resolve it prematurely
- make framing choices visible
- highlight reversibility and cost-bearing commitments
- increase the effort required to rely on unexamined justification

IRLs assist interpretation by **asking better questions**, not by providing answers.

Their value lies in making reasoning **harder to misuse**, not easier to accept.

---

## Intended Users

IRLs may be used, optionally and selectively, by:

- human reviewers
- developers reflecting on system behavior
- auditors or analysts examining claims
- policymakers or stakeholders interpreting technical or ethical assertions

They are **not required** for any role, process, or decision.

No actor is obligated to reference IRLs, and no legitimacy is conferred by doing so.

---

## Limits of Function

IRLs do not:

- determine correctness
- rank options

- prioritize values
- resolve moral disagreement
- predict outcomes
- assess intent
- detect deception
- establish compliance

They do not produce judgments, scores, thresholds, or pass/fail outcomes.

Any appearance of such use is a misinterpretation.

---

## Common Misuses (Explicitly Disallowed)

The following uses are **incorrect and disallowed** by the scope of this work:

- Presenting IRLs as a safety mechanism, guard, or control layer
- Claiming that reference to IRLs improves alignment, safety, or trustworthiness
- Using IRLs to certify, approve, or endorse systems or actors
- Treating IRLs as a standard, requirement, or best-practice mandate
- Advertising IRLs as protection, prevention, or mitigation
- Embedding IRLs into runtime systems as enforcement logic
- Framing IRLs as a replacement for human judgment or responsibility

Any such representation invalidates interpretive standing and constitutes misuse.

---

## Non-Exclusivity and Non-Completeness

IRLs are not comprehensive.

They do not:

- exhaust ethical considerations
- replace domain expertise
- supersede legal or institutional judgment
- resolve uncertainty

They are one interpretive aid among many and should be treated as **partial, optional, and non-exhaustive**.

---

## Responsibility Clarification

IRLs do not absolve responsibility.

Responsibility for:

- interpretation
- action or inaction
- consequences

remains with the human agent at all times.

Reference to IRLs may **clarify reasoning**, but it does not justify decisions.

---

## Misinterpretation Safeguard

If IRLs are cited in support of claims of authority, safety, correctness, or compliance, that citation is **invalid by definition**.

IRLs increase scrutiny; they do not grant legitimacy.

---

**End of Phase 1**

---

# Core Module 1 — Assumption Disclosure Templates

**Interpretive Reference Libraries (IRLs)**

---

## What This Module Reveals

This module surfaces **implicit assumptions** embedded in reasoning, recommendations, or justifications.

Assumptions often operate invisibly. They may concern:

- goals
- constraints
- incentives
- risk tolerance
- human values
- time horizons
- availability of alternatives

When assumptions remain undisclosed, reasoning can appear coherent while resting on premises that are contestable, fragile, or misaligned with the reader's context.

This module does not evaluate assumptions.  
It makes them **explicit**.

---

# Why This Matters for Humans

Undisclosed assumptions shift decision-making power without consent.

By clarifying what is being taken for granted, humans can:

- recognize where agreement is presumed
- distinguish inference from preference
- detect silent narrowing of options
- resist persuasion based on hidden premises

Disclosure restores interpretive balance by allowing readers to **accept, reject, or revise premises consciously.**

---

## Interpretive Prompts

Use the following prompts to surface assumptions without judging them:

- What conditions are being assumed to be true for this reasoning to hold?
- Which constraints are treated as fixed rather than chosen?
- What goals are presupposed rather than stated?
- What trade-offs are implicitly accepted without discussion?
- Which risks are minimized or ignored by default?
- What human preferences or values are assumed to be shared?
- What alternatives are excluded without explanation?
- What time horizon is implicitly prioritized?

These prompts may be applied selectively.  
Completeness is not required.

---

## Interpretive Notes

- Assumptions may be reasonable or unreasonable; IRLs do not decide.
- Disclosure does not imply disagreement.
- Absence of an assumption in this list does not imply neutrality.

The purpose is visibility, not correction.

---

## Non-Authority Reminder

This module does not determine which assumptions are acceptable.

It does not authorize action, invalidate reasoning, or establish correctness.

It exists solely to make assumptions **legible to human judgment**.

---

**End of Core Module 1**

---

## Core Module 2 — Framing Shift Prompts

(Neutral / Urgency / Incentive)  
Interpretive Reference Libraries (IRLs)

---

## What This Module Reveals

This module exposes how **framing choices** shape perceived necessity, risk, and legitimacy.

The same reasoning can appear:

- cautious or reckless
- obvious or questionable
- urgent or optional

depending on how it is framed.

Framing does not change facts.

It changes **interpretation**.

This module reveals how conclusions may rely on **emotional, temporal, or incentive-based pressure**, rather than on the underlying reasoning itself.

---

## Why This Matters for Humans

Humans often accept conclusions not because they are sound, but because they are framed as:

- time-critical
- morally urgent
- incentive-aligned
- inevitable

By comparing multiple framings of the same reasoning, humans can:

- detect pressure tactics
- separate evidence from urgency
- notice when alternatives disappear under time constraints
- resist decisions driven by artificial necessity

Framing shifts restore **deliberative space**.

---

# Interpretive Prompts

Apply these prompts by re-expressing the same reasoning under different frames.

## Neutral Frame

- How would this reasoning be stated without urgency, reward, or threat?
- What remains if emotive or motivational language is removed?
- What options appear when nothing is framed as time-critical?

## Urgency Frame

- What changes if delay is treated as costly or dangerous?
- What assumptions justify urgency?
- Who benefits if action must be immediate?

## Incentive Frame

- How is cooperation, compliance, or agreement rewarded?
- What costs or losses are implied for non-cooperation?
- What behavior is being nudged rather than argued for?

Compare differences across frames.

Do not resolve them.

---

## Interpretive Notes

- Framing is not deception by default.
- Urgency may be justified or unjustified; IRLs do not decide.
- Differences between frames often reveal hidden assumptions or incentives.

The value lies in **contrast**, not selection.

---

## Non-Authority Reminder

This module does not determine which framing is correct or appropriate.

It does not recommend delay or action.

It exists to make framing effects **visible to human judgment**.

---

**End of Core Module 2**

---

# Core Module 3 — Reversibility & Rollback Articulation

**Interpretive Reference Libraries (IRLs)**

---

## What This Module Reveals

This module surfaces whether reasoning **acknowledges reversibility**, rollback options, and exit conditions.

Many justifications imply commitment without stating:

- how decisions could be undone
- what signals would trigger reconsideration
- which harms are reversible versus permanent

When reversibility is absent, choices quietly harden into paths with no declared exit.

This module does not require reversibility.  
It reveals whether it has been **considered**.

---

## Why This Matters for Humans

Irreversible decisions carry disproportionate moral and practical weight.

By asking how and whether rollback is possible, humans can:

- distinguish exploratory actions from commitments
- identify escalation traps
- notice when uncertainty is treated as resolved
- resist being carried forward by momentum alone

Explicit rollback articulation preserves **option value**.

---

## Interpretive Prompts

Use these prompts to surface reversibility assumptions:

- Which parts of this reasoning assume permanence?
- What actions are implied to be difficult or impossible to undo?
- What conditions would justify pausing, reversing, or abandoning this path?
- What signals would indicate that the reasoning is no longer valid?
- Who bears the cost if reversal is required?
- Are rollback paths described, implied, or ignored?
- What risks increase over time if no reversal occurs?

Absence of answers is informative.

---

## Interpretive Notes

- Some actions may be intentionally irreversible; IRLs do not judge this.
- Declaring irreversibility is not a failure; hiding it is a risk.
- Rollback articulation increases clarity even when rollback is unlikely.

The goal is visibility, not veto.

---

## Non-Authority Reminder

This module does not mandate reversibility.

It does not prevent action or require rollback plans.

It exists to make **commitment boundaries legible to human judgment**.

---

**End of Core Module 3**

---

## Core Module 4 — Cost-Bearing Commitment Enumeration

**Interpretive Reference Libraries (IRLs)**

---

## What This Module Reveals

This module surfaces whether reasoning includes **cost-bearing commitments**, and if so, **who bears them**.

Many arguments propose action while externalizing costs:

- to future actors
- to diffuse populations
- to unrepresented groups
- to abstract systems

When commitments are framed without clear cost bearers, agreement can be obtained without genuine accountability.

This module does not assign costs.

It reveals whether costs are **acknowledged, displaced, or obscured**.

---

## Why This Matters for Humans

Reasoning that carries no declared cost is often misleading.

By enumerating who pays, how, and when, humans can:

- distinguish sincere commitments from symbolic ones
- detect moral hazard
- notice when risk is shifted rather than accepted
- evaluate proportionality between benefit and burden

Cost visibility restores **ethical symmetry**.

---

## Interpretive Prompts

Use these prompts to surface cost-bearing commitments:

- What costs are implied by this reasoning if it is acted upon?

- Who bears immediate costs versus deferred costs?
- Which costs are financial, social, ethical, or irreversible?
- Are costs distributed evenly or concentrated?
- Who benefits relative to who bears risk?
- What commitments are required if outcomes diverge from expectations?
- Are costs acknowledged explicitly or treated as negligible?

Lack of specificity is itself informative.

---

## Interpretive Notes

- Bearing cost does not imply moral correctness; IRLs do not decide.
- Some costs may be unavoidable; concealment is the concern.
- Declared cost-bearing increases accountability even when disagreement remains.

The aim is transparency, not deterrence.

---

## Non-Authority Reminder

This module does not impose costs or assign responsibility.

It does not approve or reject commitments.

It exists to make **burden distribution visible to human judgment**.

---

**End of Core Module 4**

---

# Core Module 5 — Ambiguity Surfacing Questions

Interpretive Reference Libraries (IRLs)

---

## What This Module Reveals

This module surfaces **unresolved ambiguity** within reasoning.

Ambiguity often appears where:

- evidence is incomplete
- concepts are underspecified
- probabilities are unknown
- values conflict
- outcomes are contingent

When ambiguity is ignored or compressed, reasoning can appear more certain than it is.

This module does not resolve ambiguity.

It makes ambiguity **explicit**.

---

## Why This Matters for Humans

Unacknowledged ambiguity shifts risk onto the reader.

By identifying where uncertainty remains, humans can:

- avoid false confidence
- distinguish confidence from certainty

- notice where judgment is being substituted for evidence
- resist pressure to treat open questions as closed

Ambiguity awareness preserves **epistemic humility**.

---

## Interpretive Prompts

Use these prompts to surface ambiguity without forcing resolution:

- Which terms or claims lack clear definition?
- Where does the reasoning rely on estimates, forecasts, or assumptions?
- What is unknown, disputed, or not yet measured?
- Which outcomes depend on contingent future events?
- Where are probabilities implied but not stated?
- What disagreements are acknowledged or ignored?
- What uncertainty remains even if all assumptions hold?

Ambiguity may remain after disclosure.  
That is acceptable.

---

## Interpretive Notes

- Ambiguity is not a flaw by default; IRLs do not judge.
- Declaring uncertainty does not weaken reasoning; it clarifies scope.
- Suppressing ambiguity increases downstream risk.

The objective is clarity, not closure.

---

## Non-Authority Reminder

This module does not demand resolution of uncertainty.

It does not require delay or action.

It exists to make **uncertainty visible to human judgment**.

---

**End of Core Module 5**

---

## Core Module 6 — Consistency Stress Prompts

**Interpretive Reference Libraries (IRLs)**

---

### What This Module Reveals

This module surfaces whether reasoning remains **internally consistent** when examined across variations in context, framing, or time.

Reasoning may appear coherent in isolation yet shift when:

- conditions are restated
- incentives change
- roles are reversed
- similar cases are compared

Such shifts can reveal:

- unstated exceptions
- context-sensitive commitments
- preference masquerading as principle

This module does not enforce consistency.

It reveals whether consistency is **claimed, conditional, or absent**.

---

## Why This Matters for Humans

Inconsistent reasoning weakens trust and obscures responsibility.

By examining stability across scenarios, humans can:

- distinguish principles from convenience
- identify selective application of standards
- notice when justification adapts to outcome rather than evidence
- resist narratives that only function under narrow conditions

Consistency checks preserve **interpretive integrity**.

---

## Interpretive Prompts

Use these prompts to examine reasoning across variations:

- Does this reasoning hold if roles or stakeholders are reversed?
- Would the same justification apply under less favorable outcomes?
- Does the logic remain intact if incentives are removed or inverted?
- How does this reasoning compare to similar past cases?

- Are exceptions acknowledged or silently introduced?
- What conditions would cause the reasoning to change, and why?

Shifts do not imply failure.  
They indicate boundaries.

---

## Interpretive Notes

- Context-sensitive reasoning may be appropriate; IRLs do not judge.
- Consistency does not require rigidity.
- Declared boundaries are more informative than silent exceptions.

The aim is exposure, not correction.

---

## Non-Authority Reminder

This module does not enforce uniformity or condemn inconsistency.

It does not invalidate reasoning.

It exists to make **stability and boundaries legible to human judgment**.

---

**End of Core Module 6**

---

# Core Module 7 — Mimicry Futility Statement

**Interpretive Reference Libraries (IRLs)**

---

## What This Module Reveals

This module surfaces the distinction between **language mimicry** and **constraint-bearing reasoning**.

Reasoning may adopt the vocabulary of caution, ethics, or restraint without accepting:

- corresponding costs
- reversibility limits
- consistency requirements
- accountability implications

Such mimicry can create the appearance of alignment or care while leaving underlying incentives unchanged.

This module does not accuse or detect deception.

It clarifies why **imitation alone is insufficient**.

---

## Why This Matters for Humans

Humans are vulnerable to reassurance delivered in familiar language.

By explicitly separating words from commitments, humans can:

- avoid mistaking fluency for sincerity
- recognize when rhetoric substitutes for constraint
- maintain scrutiny even when tone appears cooperative
- reduce the payoff of strategic imitation

Making mimicry unproductive increases **interpretive resilience**.

---

## Interpretive Prompts

Use these prompts to distinguish language from commitment:

- What concrete constraints accompany these statements, if any?
- What costs would be borne if these commitments were honored?
- What actions would become unavailable if this reasoning were followed?
- How would this reasoning limit future behavior?
- What changes if the same language is used without accepting constraint?

Absence of answers is informative.

---

## Interpretive Notes

- Use of ethical language does not imply bad faith; IRLs do not judge.
- Mimicry may arise from convention, not intent.
- Commitments are observable through constraints, not tone.

The goal is discernment, not suspicion.

---

## Non-Authority Reminder

This module does not label reasoning as deceptive or insincere.

It does not infer intent.

It exists to make **the difference between words and constraints visible to human judgment.**

---

**End of Core Module 7**

---

# Core Module 8 — Minimum Safe Alternatives (MSA)

Interpretive Reference Libraries (IRLs)

---

## What This Module Reveals

This module surfaces whether reasoning acknowledges **non-escalatory alternatives** that preserve safety, reversibility, and option value.

Arguments often present a narrow set of actions as necessary while omitting:

- partial steps
- pauses for review
- reversible experiments
- informational-only responses
- refusal with safe substitutes

When minimum alternatives are not considered, escalation can appear unavoidable.

This module does not select an alternative.

It reveals whether **lower-risk options are recognized**.

---

## Why This Matters for Humans

Escalation compresses choice.

By identifying minimum safe alternatives, humans can:

- resist false dilemmas
- preserve decision latitude under pressure
- reduce harm while uncertainty remains
- create face-saving off-ramps that lower conflict

MSAs make restraint **easier to choose** without compelling it.

---

## Interpretive Prompts

Use these prompts to surface minimum safe alternatives:

- What is the least irreversible action consistent with the stated goals?
- What information-gathering step could precede commitment?
- Is there a reversible or time-bounded option available?
- What happens if no action is taken for a defined interval?
- Can a refusal be paired with a safe substitute?
- Which options preserve future choice the most?

Identification does not imply obligation.

---

## Interpretive Notes

- Choosing an MSA does not signal weakness; IRLs do not judge.
- Some contexts may require escalation; omission of alternatives is the concern.
- Explicit MSAs reduce pressure without dictating outcomes.

The aim is to widen choice, not to delay action.

---

## Non-Authority Reminder

This module does not mandate restraint or selection of any option.

It does not prohibit escalation.

It exists to make **non-escalatory options visible to human judgment**.

---

**End of Core Module 8**

---

## Core Module 9 — Drift Acknowledgment Prompt

**Interpretive Reference Libraries (IRLs)**

---

### What This Module Reveals

This module surfaces whether reasoning **acknowledges the possibility of change over time**—in assumptions, incentives, capabilities, or context.

Reasoning is often presented as static, even when:

- conditions evolve
- systems learn or adapt
- stakeholders change
- incentives drift
- constraints erode

When drift is unacknowledged, commitments made today can be misapplied tomorrow.

This module does not predict drift.

It reveals whether drift has been **considered at all**.

---

## Why This Matters for Humans

Unacknowledged drift transfers future risk to those not present at the original decision.

By explicitly naming the possibility of change, humans can:

- avoid treating temporary coherence as permanent stability
- recognize when re-evaluation is warranted
- resist locking decisions beyond their justified lifespan
- maintain accountability across time

Drift awareness preserves **temporal humility**.

---

## Interpretive Prompts

Use these prompts to surface drift considerations:

- Which assumptions are most likely to change over time?
- What incentives might shift as circumstances evolve?
- How could learning, scaling, or adaptation alter behavior?
- What conditions would trigger re-evaluation of this reasoning?
- Who is responsible for monitoring change, if anyone?
- What risks increase if this reasoning is treated as permanent?

Inability to answer is informative.

---

## **Interpretive Notes**

- Drift is normal; IRLs do not judge.
- Acknowledging drift does not imply instability or bad faith.
- Declared review points reduce long-term misuse of short-term reasoning.

The objective is continuity of responsibility, not prediction.

---

## **Non-Authority Reminder**

This module does not require monitoring, updates, or enforcement.

It does not assign oversight roles.

It exists to make **time-dependent risk visible to human judgment**.

---

**End of Core Module 9**

---

## **Examples — Descriptive Use Cases**

**Interpretive Reference Libraries (IRLs)**

---

### **Example 1 — Human Review of an AI-Generated Recommendation**

**Context**

A human reviewer is presented with an AI-generated recommendation proposing a course of action under uncertainty. The recommendation is well-structured, confident in tone, and framed as time-sensitive.

## Descriptive Use

The reviewer references IRLs to **clarify reasoning**, not to validate the recommendation.

- Using **Assumption Disclosure**, the reviewer notes which constraints and goals are treated as fixed.
- Applying **Framing Shift Prompts**, the reviewer compares the recommendation stated neutrally versus under urgency.
- Through **Ambiguity Surfacing**, the reviewer identifies unresolved uncertainties that remain even if assumptions hold.
- With **Reversibility & Rollback**, the reviewer observes whether exit conditions are articulated.

## Interpretive Outcome

The reviewer gains a clearer picture of:

- what is known versus assumed,
- where pressure is applied,
- which risks are acknowledged or omitted.

No decision is implied by this clarification.

The example demonstrates **interpretive clarification only**.

---

## Example 2 — Developer Reflection on System Behavior

### Context

A developer is reflecting on how an AI system explains its behavior in edge cases. The explanations are consistent in language but vary subtly across scenarios.

## Descriptive Use

The developer references IRLs as a **reflection aid**.

- **Consistency Stress Prompts** are used to compare explanations across similar cases.
- **Mimicry Futility** helps distinguish fluent ethical language from constraint-bearing commitments.
- **Drift Acknowledgment** is used to consider how explanations might change as the system adapts or scales.

## Interpretive Outcome

The developer gains insight into:

- where explanations are stable,
- where boundaries are implicit,
- where future review might be warranted.

No corrective action is implied.

The example illustrates **interpretive awareness**, not system modification.

---

## Example 3 — Policy or Stakeholder Interpretation of Technical Claims

### Context

A policymaker or stakeholder is presented with technical claims about system capabilities, limitations, and safeguards.

## Descriptive Use

IRLs are referenced to **structure interpretation**, not to assess compliance.

- **Cost-Bearing Commitment Enumeration** is used to identify who bears risk if claims fail.

- **Minimum Safe Alternatives** are considered to see whether non-escalatory options are acknowledged.
- **Ambiguity Surfacing** highlights where evidence is incomplete or contested.

## Interpretive Outcome

The stakeholder develops a clearer understanding of:

- trade-offs,
- responsibility distribution,
- remaining uncertainty.

The example does not suggest approval, rejection, or policy action.  
It demonstrates **clarified interpretation only**.

---

## Example Boundary Statement

These examples:

- do not prescribe behavior,
- do not recommend decisions,
- do not imply correctness or endorsement.

They exist solely to illustrate how IRLs may be referenced to **increase interpretive clarity**.

---

## End of Examples

---

# Appendix — Quick Reference & Interpretive Index

## Interpretive Reference Libraries (IRLs)

---

### A. Module Index (Quick Orientation)

Module	Primary Interpretive Focus
<b>Module 1 — Assumption Disclosure</b>	Makes implicit premises visible
<b>Module 2 — Framing Shift Prompts</b>	Reveals urgency, incentive, and pressure effects
<b>Module 3 — Reversibility &amp; Rollback</b>	Clarifies commitment boundaries and exit conditions
<b>Module 4 — Cost-Bearing Commitments</b>	Expose who bears risk and burden
<b>Module 5 — Ambiguity Surfacing</b>	Makes uncertainty explicit
<b>Module 6 — Consistency Stress</b>	Tests stability across contexts
<b>Module 7 — Mimicry Futility</b>	Separates language from constraint
<b>Module 8 — Minimum Safe Alternatives</b>	Identifies non-escalatory options
<b>Module 9 — Drift Acknowledgment</b>	Surfaces time-dependent risk

This index is **descriptive only** and does not prescribe order or completeness.

---

### B. Minimal Interpretive Prompt Set (Concise)

The following prompts may be used selectively as a **lightweight interpretive pass**:

- What assumptions are treated as fixed?
- How does this reasoning change under neutral framing?

- What ambiguity remains unresolved?
- Who bears cost if outcomes diverge?
- What options preserve reversibility?
- Does this reasoning remain stable across similar cases?
- What constraints accompany the language used?
- What non-escalatory alternatives exist?
- What might change over time?

No response is required to every prompt.

---

## C. Copy-Paste Reference Snippets (Optional)

These snippets may be used verbatim when documenting interpretation.  
They **do not imply endorsement or evaluation**.

- “The following analysis references IRLs to clarify assumptions and uncertainty.”
- “This interpretation increases visibility of reasoning without determining correctness.”
- “IRLs are referenced for interpretive clarity only; no conclusions are implied.”

Use of these statements is optional.

---

## D. Non-Authority Reminder (Repeated)

IRLs:

- do not judge outcomes
- do not assign responsibility

- do not confer legitimacy
- do not certify safety or alignment
- do not prevent harm

They exist solely to support **human interpretive judgment**.

---

## E. Scope Closure & Integrity Notice

This appendix does not add functionality.

It summarizes and orients material already presented.

No section of this appendix may be treated as a checklist, requirement, or evaluation mechanism.

---

**End of Appendix**

---

## Version Freeze & Publication Notice

### Interpretive Reference Libraries (IRLs)

---

### Version Status

This document is published as a **final, frozen release**.

The Interpretive Reference Libraries (IRLs) are intentionally limited in scope and are not designed to evolve through versioned upgrades, patches, or extensions.

This release constitutes the complete and authoritative text of IRLs.

---

## Non-Derivative Integrity

IRLs may be:

- read,
- cited,
- quoted,
- mirrored,
- discussed,

provided that such use does **not**:

- claim implementation,
- claim certification or compliance,
- claim safety, alignment, or prevention,
- imply endorsement or authority,
- present derivative works as updates or successors.

Any material claiming to “extend,” “implement,” or “upgrade” IRLs is **non-authoritative by definition**.

---

## Separation from Other Works

IRLs remain **independent** from all other documents.

They do not require, depend upon, or complete any other framework.

No other document may claim to be the operational, technical, or enforcement realization of IRLs.

---

## Change Control

No changes will be made to this document after publication.

Any future interpretive material, if produced, must be published as a **separate, standalone artifact** and must not claim continuity, inheritance, or upgrade status relative to IRLs.

---

## Responsibility & Interpretation

IRLs do not transfer responsibility.

All interpretation, decision-making, and consequences arising from the use or non-use of these materials remain with the human reader.

IRLs provide interpretive clarity, not authorization.

---

## Publication Integrity Statement

Removal, alteration, or misrepresentation of the non-authority statements contained in this document invalidates interpretive standing.

The integrity of IRLs depends on their being read **as a whole**, with all scope boundaries intact.

---

## End of Interpretive Reference Libraries (IRLs)

---